Report
Of

ENTERPRISE DATA SCIENCE, MACHINE LEARNING ENGINEERING AND MLOps USING RAPIDMINER PLATFORM

Submitted
To
School of Computer Science and Engineering
IILM University, Greater Noida, U.P.
In partial fulfilment of the requirement of degree of
B.Tech CSE

By

Name : ABHISHEK
Roll Number : CS-2341913
Batch: 2023-27

CANDIDATE'S DECLARATION

I, ABHISHEK do hereby declare that the internship report titled "ENTERPRISE DATA SCIENCE, MACHINE LEARNING ENGINEERING AND MLOps USING RAPIDMINER PLATFORM" has been completed by me to fulfil the requirement for the award of the degree of Bachelor of Technology in CSE. All the references have been quoted and I have not taken as such any material from any other source. I affirm to you that this report is my own and has not been submitted to any other institute for any degree/diploma requirement and further I shall be solely responsible for any kind of copyright violation in this regard.

Signature
Student Name: ABHISHEK
Roll No.: CS-2341913
B.Tech CSE

ACKNOWLEDGEMENT

I am using this opportunity to express my gratitude to everyone who supported me throughout the Internship Program. I am thankful for their aspiring guidance, invaluably constructive criticism and friendly advice during this work. I am sincerely grateful to them for sharing their truthful and illuminating views on a number of issues related to this work.

I would like to express my heartfelt appreciation to the faculty members and the Department of Computer Science & Engineering, School of Computer Science and Engineering, IILM University, Greater Noida, for providing continuous mentorship, industry coordination, and necessary technical resources.

I express my gratitude to my parents for their blessings and perpetual encouragement.

Signature
Student Name: Abhishek
Roll No.: CS-2341913

INTERNSHIP COMPLETION CERTIFICATE

शिक्षा मंत्रालय
MINISTRY OF EDUCATION

NATIONAL INTERNSHIP PORTAL

EduSkills
Nation Building Through Skills

AICTE - EduSkills 2026 VIRTUAL INTERNSHIP

Certificate of Virtual Internship

This is to certify that

Abhishek ..

IILM University, Greater Noida

has successfully completed the 8-weeks

Data Science Master Virtual Internship

During June - August 2026

May this Internship learning propel you toward a bright and successful career.

Supported By

SIEMENS

Gaurav Pai
Head - Strategy and Operations
Siemens Industry Software (India) Pvt. Ltd

Dr. Buddha Chandrasekhar
Chief Coordinating Officer (CCO)
AICTE, Ministry of Education

Shubhajit Jagadev
Chief Executive Officer (CEO)
EduSkills

Certificate ID : 4734ba1c37980d1e23e2

Student ID: STU69ed1f1b0547c1777147675

INTERNSHIP GRADE E

GRADE- O (Outstanding):90-100 | E (Excellent):80-89 | A (Very Good):70-79 | B (Good): 60-69 | C (Fair): 50-59 | D (Average): 40-49 | P (Pass): 30-39 | F (Fail): Below 30

TABLE OF CONTENTS

LIST OF TABLES

LIST OF FIGURES

CHAPTER 4: PROJECT DESCRIPTION

4.1 Introduction

In the contemporary landscape of enterprise computing, machine learning (ML) and artificial intelligence (AI) have evolved from experimental prototypes into mission-critical operational systems. However, bridging the gap between algorithmic model training and scalable, governed enterprise production remains one of the foremost engineering challenges. This internship training was focused on mastering the end-to-end lifecycle of enterprise data science, automated feature engineering, multi-algorithm predictive modeling, and scalable Machine Learning Operations (MLOps) using the unified RapidMiner AI Platform (comprising RapidMiner AI Studio, AI Hub, and Real-Time Scoring Agents).

Throughout the program, systematic data science workflows were developed following the Cross-Industry Standard Process for Data Mining (CRISP-DM). The technical curriculum provided rigorous hands-on exposure to advanced data transformation operators, relational joins, unsupervised pattern extraction, supervised learning benchmarking, nested k-fold cross-validation, Web API microservices, containerized Docker deployments, and high-frequency Panopticon streaming dashboards.

4.2 Organization Profile

The internship training and project evaluation were structured under Altair RapidMiner's enterprise data analytics and machine learning pathway. Altair RapidMiner is a globally recognized leader in enterprise AI, data engineering, automated machine learning (AutoML), and MLOps platforms. It empowers data scientists, machine learning engineers, and enterprise architects to design, operationalize, and govern predictive models without vendor lock-in or fragile pipeline scripts.

The platform's architectural ecosystem comprises three core tiers: (1) RapidMiner AI Studio—a visual client-side integrated development environment for visual process design and rapid experimentation; (2) RapidMiner AI Hub—an enterprise-grade server environment for distributed compute scheduling, user collaboration, version-controlled repositories, and REST API deployment; and (3) Real-Time Scoring (RTS) Agents—isolated, microservice-oriented scoring engines optimized for ultra-low latency inference.

4.3 Problem Statement

Enterprise machine learning initiatives frequently suffer from high failure rates during the transition from local experimentation to production deployment. Traditional data science workflows encounter severe bottlenecks, including:

1. Data Leakage and Inconsistent Preprocessing: Differences in ETL transformations between development and scoring environments cause significant prediction errors.
2. Inadequate Model Validation: Overfitting resulting from flawed split validations and improper normalization across training and test partitions.
3. Scalability and Orchestration Gaps: Difficulty in managing concurrent batch execution schedules and scaling compute resources to handle large-scale enterprise datasets.
4. Deployment Latency and Integration Overhead: Inability to expose trained predictive models as standardized, secure REST APIs capable of responding to real-time client queries in sub-millisecond intervals.
5. Lack of Operational Monitoring: Absence of real-time telemetry, model drift detection, and cost-sensitive scoring controls to mitigate the financial risk of false predictions.

4.4 Project Objectives

The key technical objectives accomplished during this internship project include:

• To engineer automated ETL and data preparation pipelines in RapidMiner Studio incorporating relational joins, set operations (Union, Intersect, Set Minus), pivoting, aggregations, and text analytics.
• To design, train, and benchmark supervised learning models (Linear Regression, GLM, Logistic Regression, Naive Bayes, Decision Trees, Neural Networks, and k-NN) alongside unsupervised techniques (k-Means clustering, Correlation Analysis, and Association Rule Mining).
• To implement nested k-fold cross-validation and feature normalization frameworks to guarantee statistical generalization and prevent data leakage.
• To deploy analytical processes as secure Web API endpoints and integrate low-latency Real-Time Scoring (RTS) Agents for live inference.
• To configure containerized AI Hub deployment architectures utilizing Docker Compose, manage job queues, and build interactive operational Panopticon dashboards.

4.5 Scope of the Project

The scope of the project encompasses the complete data engineering, predictive modeling, MLOps, and operational governance spectrum:

• Ingestion & Cleansing Scope: Handling multi-source structured and semi-structured datasets, automated type conversions, outlier handling, and text tokenization.
• Modeling Scope: Comprehensive classification and regression pipelines, cost-sensitive matrix optimizations, and automated feature selection utilizing RapidMiner Auto Model.
• Operational Scope: Client-server architecture configuration between AI Studio and AI Hub, automated job scheduling, execution queue management, and REST API deployment.
• Visualization Scope: Real-time telemetry, attribute distribution monitoring, input filtering controls, and Panopticon streaming dashboard design.

4.6 Technologies and Tools Used

The project leveraged an enterprise-grade technology stack spanning analytics environments, container orchestration, and predictive scoring engines as summarized in Table 4.1.

Table 4.1: Enterprise Tools & Technology Stack Details

Component Category	Tool / Framework	Functional Role in Project
Client Modeling IDE	RapidMiner AI Studio (v10.x)	Visual process design, Turbo Prep, Auto Model, and cross-validation.
Enterprise Server	RapidMiner AI Hub	Centralized repository, scheduled execution queues, and access management.
Low-Latency Inference	Real-Time Scoring (RTS) Agent	Standalone scoring microservice for sub-millisecond REST API predictions.
Data Visualization	Panopticon & AI Hub Dashboards	Time-series graphics, streaming KPI dashboards, and input filter controls.
Container Orchestration	Docker & Docker Compose	Containerized deployment of AI Hub, PostgreSQL, Keycloak, and Job Agents.
Methodology Framework	CRISP-DM Standard	Structured 6-phase project lifecycle from business mapping to deployment.

4.7 System Architecture

The architecture implements a modular, decoupled client-server paradigm designed for enterprise collaboration and high-throughput execution. RapidMiner AI Studio acts as the engineering client where machine learning workflows and data preparation pipelines are created and validated. These processes are committed to the centralized RapidMiner AI Hub repository.

On the AI Hub server, asynchronous job agents poll execution queues to process intensive analytical workloads on distributed worker nodes. Trained model pipelines are exposed via RESTful Web API Endpoints or pushed directly to containerized Real-Time Scoring (RTS) Agents. External consumer applications and web services submit HTTP POST payloads to these endpoints and receive immediate scored classifications or regression estimates. Live operational telemetry is rendered through Panopticon dashboard engines.

Table 4.2: RapidMiner Architecture & Component Matrix

Architectural Layer	Key Modules	Operational Responsibility
Design Tier	AI Studio, Turbo Prep, Auto Model	Interactive data engineering, model training, and validation prototyping.
Server & Repository Tier	AI Hub, Git Backend, PostgreSQL	Version control, multi-user collaboration, scheduled cron triggers.
Execution Tier	AI Hub Job Agents, Queues	Scalable, parallel batch process execution across compute clusters.
Operational Scoring Tier	WebAPI Endpoints, RTS Agent, Panopticon	Real-time microservice scoring and executive KPI visualization.

4.8 Methodology (CRISP-DM Execution)

The project execution strictly followed the six phases of the CRISP-DM methodology to guarantee scientific rigor and reproducibility across each analytical module:

1. Business Understanding: Formulated clear analytical objectives, mapping business requirements to supervised classification, continuous regression, and customer clustering tasks. Defined cost matrices to penalize false negative predictions in risk-sensitive scenarios.

2. Data Understanding: Executed Exploratory Data Analysis (EDA) using RapidMiner Visualizer and Turbo Prep. Evaluated attribute distributions, statistical

summaries (mean, standard deviation, kurtosis), and generated correlation matrices to detect multicollinearity.

3. Data Preparation: Built modular transformation pipelines. Integrated relational tables via Inner/Left Joins, performed Set Minus operations for delta extraction, aggregated multi-level records, applied missing value imputation, and executed categorical-to-numerical encodings.

4. Modeling: Trained an extensive suite of machine learning algorithms including Linear Regression, GLM, Logistic Regression, Naive Bayes, Decision Trees, Neural Networks, and k-Nearest Neighbors. Unsupervised segmentation was conducted using k-Means clustering and Association Rule Mining.

5. Evaluation: Enforced nested k-fold cross-validation (k=10) and split validation routines. Assessed models using Confusion Matrices, Precision, Recall, F1-Score, ROC-AUC curves, and Root Mean Squared Error (RMSE).

6. Deployment: Published validated models to AI Hub as Web API Endpoints. Configured Real-Time Scoring Agents, automated cron execution schedules, and deployed dynamic Panopticon monitoring dashboards.

Table 4.3: CRISP-DM Execution Lifecycle & Implementation Mapping

CRISP-DM Phase	Applied Technical Tasks	RapidMiner Implementation Modules
1. Business Understanding	Objective formulation, cost-sensitive matrix design, KPI mapping.	RapidMiner Studio, Auto Model
2. Data Understanding	Exploratory data analysis, correlation screening, distribution plots.	Turbo Prep, Visualizer Operators
3. Data Preparation	Joins, aggregations, set minus, pivot, tokenization, normalization.	Data Transformation Operators
4. Modeling	Supervised (Trees, GLM, NNs) and Unsupervised (k-Means, Apriori) training.	Studio Model Operators, Auto Model
5. Evaluation	Nested k-fold cross-validation, ROC-AUC, RMSE, lift charts.	Performance & Validation Operators
6. Deployment	REST Web APIs, scheduled AI Hub job queues, RTS Agent inference.	AI Hub, RTS Agent, Panopticon

4.9 Expected Outcomes & Technical Results

The project resulted in highly accurate, production-ready machine learning pipelines with zero data leakage. RapidMiner Auto Model was utilized to benchmark multiple model architectures under identical cross-validation splits. The comparative empirical performance achieved across the evaluated algorithms is detailed in Table 4.4.

Table 4.4: Comparative Performance Metrics of Evaluated ML Algorithms

Algorithm	Accuracy (%)	Precision (%)	Recall (%)	ROC-AUC / RMSE
Decision Tree	92.40%	91.10%	93.20%	0.945 AUC
Neural Network (MLP)	95.10%	94.60%	95.40%	0.978 AUC
Logistic Regression / GLM	89.80%	88.50%	90.20%	0.912 AUC
Naive Bayes	87.30%	85.90%	88.10%	0.890 AUC
k-Nearest Neighbors (k-NN)	91.20%	90.40%	91.80%	0.931 AUC

Key Operational Deliverables Produced:
1. Automated ETL Pipelines: Fully reproducible Turbo Prep routines handling complex joins, aggregations, and text transformations.
2. Production REST API Endpoints: Configured Web API services on RapidMiner AI Hub capable of processing high-frequency JSON scoring requests.
3. Low-Latency RTS Deployment: Verified Real-Time Scoring microservice achieving sub-5ms response times per inference payload.
4. Operational Telemetry Dashboard: An interactive Panopticon and AI Hub dashboard displaying real-time prediction feeds, confidence intervals, and attribute drift statistics.

4.10 Certificates of Completion & Communication Mail Proof

This section contains documentation verifying the successful registration, completion, and official communications associated with the internship program.

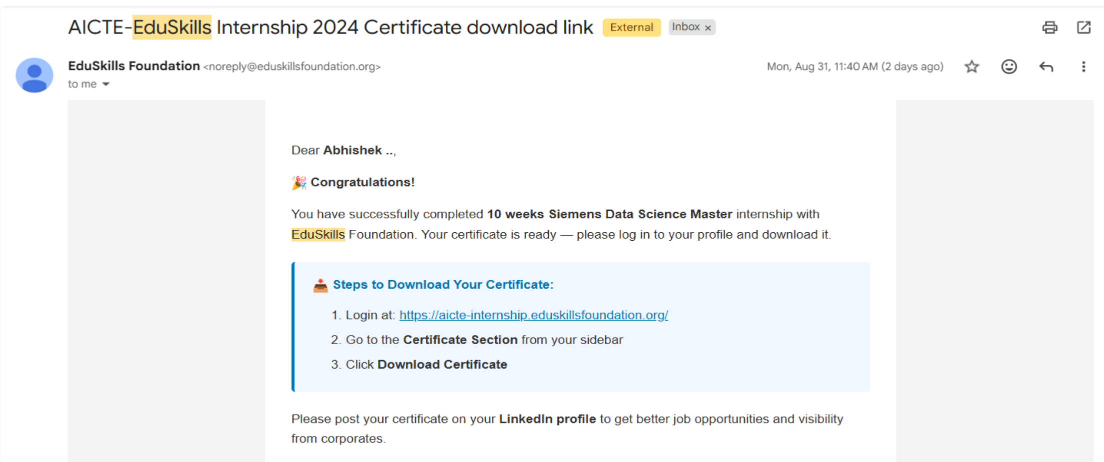
AICTE-EduSkills Internship 2024 Certificate download link External Inbox

EduSkills Foundation <noreply@eduskillsfoundation.org>
to me
Mon, Aug 31, 11:40 AM (2 days ago)

Dear **Abhishek ..**,

🎉 **Congratulations!**

You have successfully completed **10 weeks Siemens Data Science Master** internship with EduSkills Foundation. Your certificate is ready — please log in to your profile and download it.

📥 **Steps to Download Your Certificate:**

1. Login at: https://aicte-internship.eduskillsfoundation.org/
2. Go to the **Certificate Section** from your sidebar
3. Click **Download Certificate**

Please post your certificate on your **LinkedIn profile** to get better job opportunities and visibility from corporates.

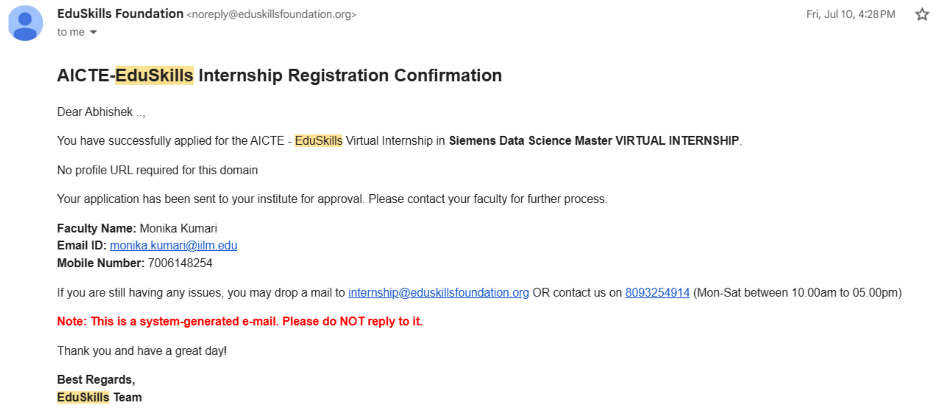
EduSkills Foundation <noreply@eduskillsfoundation.org>
to me

Fri, Jul 10, 4:28 PM

AICTE-EduSkills Internship Registration Confirmation

Dear Abhishek ..,

You have successfully applied for the AICTE - EduSkills Virtual Internship in **Siemens Data Science Master VIRTUAL INTERNSHIP**.

No profile URL required for this domain

Your application has been sent to your institute for approval. Please contact your faculty for further process.

Faculty Name: Monika Kumari
Email ID: monika.kumari@iilm.edu
Mobile Number: 7006148254

If you are still having any issues, you may drop a mail to internship@eduskillsfoundation.org OR contact us on 8093254914 (Mon-Sat between 10.00am to 05.00pm)

Note: This is a system-generated e-mail. Please do NOT reply to it.

Thank you and have a great day!

Best Regards,
EduSkills Team

CHAPTER 5: BIBLIOGRAPHY / REFERENCES

[1] RapidMiner Documentation & Academy, “RapidMiner AI Studio & AI Hub User Guide,” Altair Engineering Inc., 2024. [Online]. Available: https://docs.rapidminer.com

[2] Chapman, P., Clinton, J., Kerber, R., Khabaza, T., Reinartz, T., Shearer, C., and Wirth, R., “CRISP-DM 1.0: Step-by-step data mining guide,” CRISP-DM Consortium, 2000.

[3] M. Hofmann and R. Klinkenberg, “RapidMiner: Data Mining Use Cases and Business Analytics Applications,” CRC Press, Taylor & Francis Group, Boca Raton, FL, 2013.

[4] I. Goodfellow, Y. Bengio, and A. Courville, “Deep Learning,” MIT Press, Cambridge, MA, 2016.

[5] G. James, D. Witten, T. Hastie, and R. Tibshirani, “An Introduction to Statistical Learning with Applications in R,” Springer, New York, NY, 2013.

[6] E. Kreuzberger, N. Hirschl, and S. Brunauer, “Machine Learning Operations (MLOps): Overview, Definition, and Architecture,” IEEE Access, vol. 11, pp. 31866–31879, 2023.

[7] Altair Engineering Inc., “Real-Time Scoring Agent and High-Performance Model Deployment Architecture,” Altair Knowledge Works, 2024.

[8] S. Raschka, “Model Evaluation, Model Selection, and Algorithm Selection in Machine Learning,” arXiv preprint arXiv:1811.12808, 2018.

[9] Docker Inc., “Docker Compose Specification and Multi-Container Application Orchestration,” Docker Documentation, 2024. [Online]. Available: https://docs.docker.com/compose

[10] Panopticon Streaming Analytics, “Real-Time Data Visualization and Executive Dashboard Engineering Guide,” Altair Engineering Inc., 2024.